

Using R with Single Variables

Working with Diamonds Dataset

Loading the diamonds data set.

I noticed that the diamonds data set will be listed as a 'Promise' in the workspace. This is a special object in R, and I need to run a command on the data to fully load the data set. I need to run summary in order to get the data appear in the environments.

```
#must load the ggplot package first
library(ggplot2)

#loads the diamonds data set since it comes with the ggplot package
data(diamonds)

summary(diamonds)
```

```
##      carat      cut      color      clarity
## Min.   :0.2000 Fair      : 1610 D: 6775 SI1      :13065
## 1st Qu.:0.4000 Good      : 4906 E: 9797 VS2      :12258
## Median :0.7000 Very Good:12082 F: 9542 SI2      : 9194
## Mean   :0.7979 Premium  :13791 G:11292 VS1      : 8171
## 3rd Qu.:1.0400 Ideal      :21551 H: 8304 VVS2     : 5066
## Max.   :5.0100              I: 5422 VVS1     : 3655
##              J: 2808 (Other): 2531
##      depth      table      price      x
## Min.   :43.00 Min.   :43.00 Min.   : 326 Min.   : 0.000
## 1st Qu.:61.00 1st Qu.:56.00 1st Qu.: 950 1st Qu.: 4.710
## Median :61.80 Median :57.00 Median : 2401 Median : 5.700
## Mean   :61.75 Mean   :57.46 Mean   : 3933 Mean   : 5.731
## 3rd Qu.:62.50 3rd Qu.:59.00 3rd Qu.: 5324 3rd Qu.: 6.540
## Max.   :79.00 Max.   :95.00 Max.   :18823 Max.   :10.740
##
##      y      z
## Min.   : 0.000 Min.   : 0.000
## 1st Qu.: 4.720 1st Qu.: 2.910
## Median : 5.710 Median : 3.530
## Mean   : 5.735 Mean   : 3.539
## 3rd Qu.: 6.540 3rd Qu.: 4.040
## Max.   :58.900 Max.   :31.800
##
```

There are 53940 observations and 10 variables in the dataset. There are 3 ordered factors in the dataset: cut, color, clarity

```
str(diamonds)

## Classes 'tbl_df', 'tbl' and 'data.frame': 53940 obs. of 10 variables:
## $ carat : num 0.23 0.21 0.23 0.29 0.31 0.24 0.24 0.26 0.22 0.23 ...
## $ cut : Ord.factor w/ 5 levels "Fair"<"Good"<...: 5 4 2 4 2 3 3 3 1 3 ...
## $ color : Ord.factor w/ 7 levels "D"<"E"<"F"<"G"<...: 2 2 2 6 7 7 6 5 2 5 ...
## $ clarity: Ord.factor w/ 8 levels "I1"<"SI2"<"SI1"<...: 2 3 5 4 2 6 7 3 4 5 ...
## $ depth : num 61.5 59.8 56.9 62.4 63.3 62.8 62.3 61.9 65.1 59.4 ...
## $ table : num 55 61 65 58 58 57 57 55 61 61 ...
## $ price : int 326 326 327 334 335 336 336 337 337 338 ...
```

```
## $ x      : num  3.95 3.89 4.05 4.2 4.34 3.94 3.95 4.07 3.87 4 ...
## $ y      : num  3.98 3.84 4.07 4.23 4.35 3.96 3.98 4.11 3.78 4.05 ...
## $ z      : num  2.43 2.31 2.31 2.63 2.75 2.48 2.47 2.53 2.49 2.39 ...
```

```
levels(diamonds$color)
```

```
## [1] "D" "E" "F" "G" "H" "I" "J"
```

According to the documents, level 'D' is the best level for the diamonds.

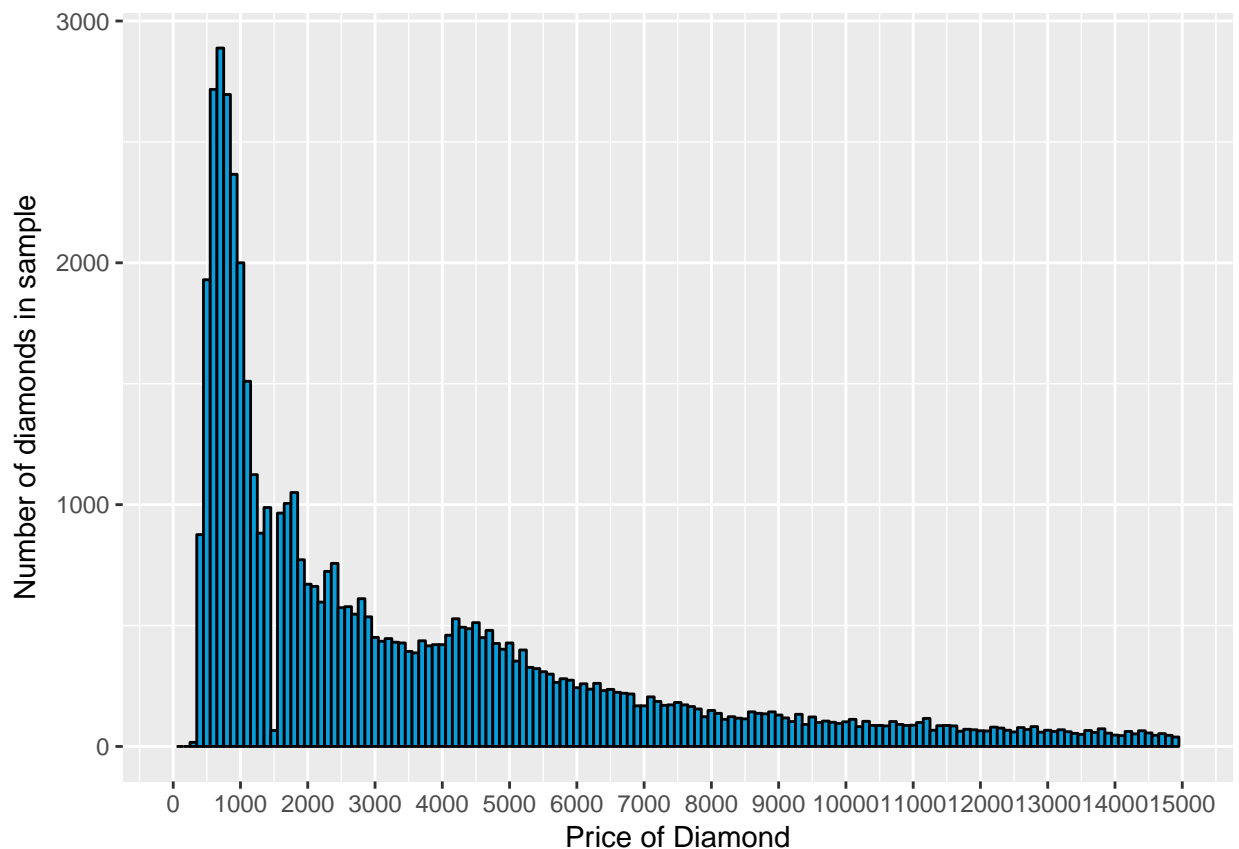
Price Histogram

Getting a histogram of the price shows a long-tailed distribution.

```
qplot(data = diamonds, x = price,
      color = I('black'), fill = I('#099DD9'),
      binwidth = 100) +
  xlab('Price of Diamond') +
  ylab('Number of diamonds in sample') +
  scale_x_log10() +
  scale_x_continuous(limits = c(0,15000), breaks = seq(0, 15000, 1000))
```

```
## Scale for 'x' is already present. Adding another scale for 'x', which
## will replace the existing scale.
```

```
## Warning: Removed 1655 rows containing non-finite values (stat_bin).
```



The shape of the price distribution is long-tailed. This means a large portion of the distribution is far from the 'head' or center of the distribution.

Running summary on price, we'll see:

```
summary(diamonds$price)

##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##      326     950    2401    3933    5324   18820
```

Now, I want to get more statistics on the price of the diamonds. I answer questions like:

How many diamond cost less than \$500? To answer this question, first I get a summary of the price for those under \$500. After executing the below command, the summary looks logical and gives true and false statements. True for those numbers below \$500 and False for those above \$500.

```
summary(diamonds$price < 500)

##      Mode  FALSE     TRUE    NA's
## logical  52211    1729         0
```

To make the above logical statement look nicer, I create a new variable called 'less_than_500' and assign it to 'NA'. Then I create a conditional statement: If the price is under \$500, return 1; else return 0.

```
less_than_500 <- NA
diamonds$less_than_500 <- ifelse(diamonds$price < 500, 1, 0)
diamonds$less_than_500 <- factor(diamonds$less_than_500)
summary(diamonds$less_than_500)
```

```
##      0      1
## 52211  1729
```

Next question, how many diamonds cost less than \$250? The same approach for the previous question will be applied here.

```
less_than_250 <- NA
diamonds$less_than_250 <- ifelse(diamonds$price < 250, 1, 0)
diamonds$less_than_250 <- factor(diamonds$less_than_250)
summary(diamonds$less_than_250)
```

```
##      0
## 53940
```

The above calculations show that the conditional statement returns only those that do NOT cost less \$250 dollars. This means there are no diamonds that cost less than \$250.

Nest question, how many diamonds cost \$15000 or more? Let's apply the same approach in the 2 previous questions to this question as well.

```
more_than_15000 <- NA
diamonds$more_than_15000 <- ifelse(diamonds$price >= 15000, 1, 0)
diamonds$more_than_15000 <- factor(diamonds$more_than_15000)
summary(diamonds$more_than_15000)
```

```
##      0      1
## 52284  1656
```

There are 1656 diamonds that cost \$15000 or more.

Cheaper Diamonds

Now, I want to explore the largest peak on the histogram I created earlier. I will use different binwidth, and limit the x-axis to observe the largest peak better.

According to my previous histogram, it's safe to say that the peak occurred somewhere between 0-2000. I decided to set limits and breaks to get a closer look on the peak. I first started to set the limit(0,2000), and breaks=seq(0,2000,100), which gave me a much clearer view that I can easily ignore anything under \$300.

Also, I could see there were no diamonds that cost \$15000.

I decided to set my final limit between \$300-1600.

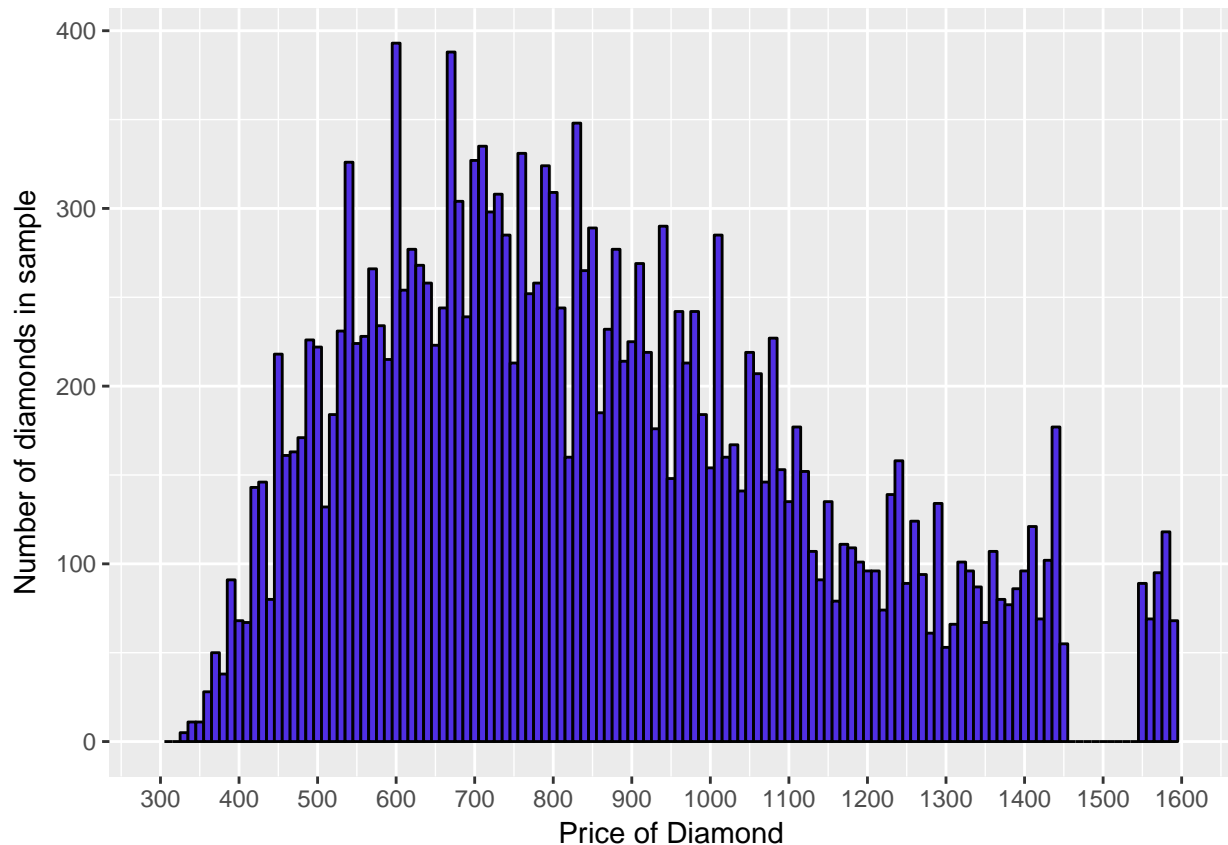
```
qplot(data = diamonds, x = price,
      color = I('black'), fill = I('#502FE4'),
      binwidth = 10) +
  xlab('Price of Diamond') +
  ylab('Number of diamonds in sample') +
  scale_x_log10() +
  scale_x_continuous(limits = c(300,1600), breaks = seq(0, 1600, 100)) +
  ggsave('priceHistogram.png')
```

```
## Scale for 'x' is already present. Adding another scale for 'x', which
## will replace the existing scale.
```

```
## Saving 6.5 x 4.5 in image
```

```
## Warning: Removed 33450 rows containing non-finite values (stat_bin).
```

```
## Warning: Removed 33450 rows containing non-finite values (stat_bin).
```



From this histogram, we can see the main peaks are between \$600-900.

Price by Cut Histograms

I want to break the histogram of diamonds prices by their cut. Let's have a look at the cut summary.

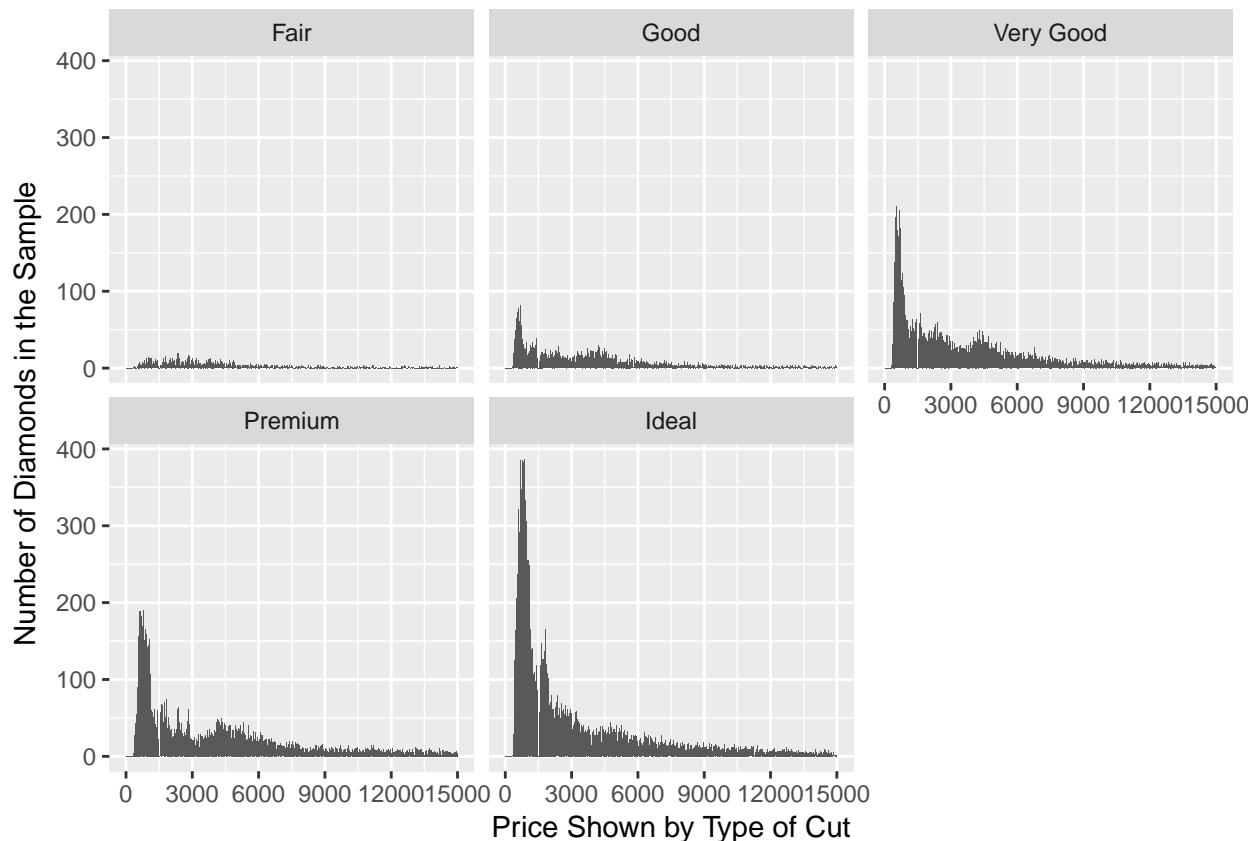
```
summary(diamonds$cut)
```

```
##      Fair      Good Very Good   Premium    Ideal  
##    1610     4906     12082    13791    21551
```

There are 5 categories of cut: Fair, Good, Very Good, Premium, and Ideal.

```
qplot(data = diamonds, x = price, binwidth = 25) +  
  xlab('Price Shown by Type of Cut') +  
  ylab('Number of Diamonds in the Sample') +  
  scale_x_continuous(limits = c(0, 15000), breaks = seq(0, 15000, 3000)) +  
  facet_wrap(~cut)
```

```
## Warning: Removed 1655 rows containing non-finite values (stat_bin).
```



Diving the prices based on the diamond cut, shows that the distribution looks similar in Ideal, Very Good and Premium categories; all long-tailed with a price peak. In the Good category, it is also long-tailed but the peak is not as much as in the first 3. In the Fair category the distribution looks uniform on the left with long tails on the right.

Price by Cut

Now, let's answer some questions about the price of the diamonds based on their cuts:

Which cut has the highest priced diamond?

To answer this question, I use 'by' to get the summary of each cut. The information below shows that Very Good and Premium cuts have the maximum value of \$18820.

```
by(diamonds$price, diamonds$cut, summary)

## diamonds$cut: Fair
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##   337   2050   3282   4359   5206   18570
## -----
## diamonds$cut: Good
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##   327   1145   3050   3929   5028   18790
## -----
## diamonds$cut: Very Good
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##   336    912   2648   3982   5373   18820
## -----
## diamonds$cut: Premium
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##   326   1046   3185   4584   6296   18820
## -----
## diamonds$cut: Ideal
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##   326    878   1810   3458   4678   18810
```

Which cut has the lowest priced diamond?

Cuts Ideal and Premium have the lowest priced diamond \$326.

Which cut has the lower median price?

Cut Ideal has the lowest Median price of \$1810.

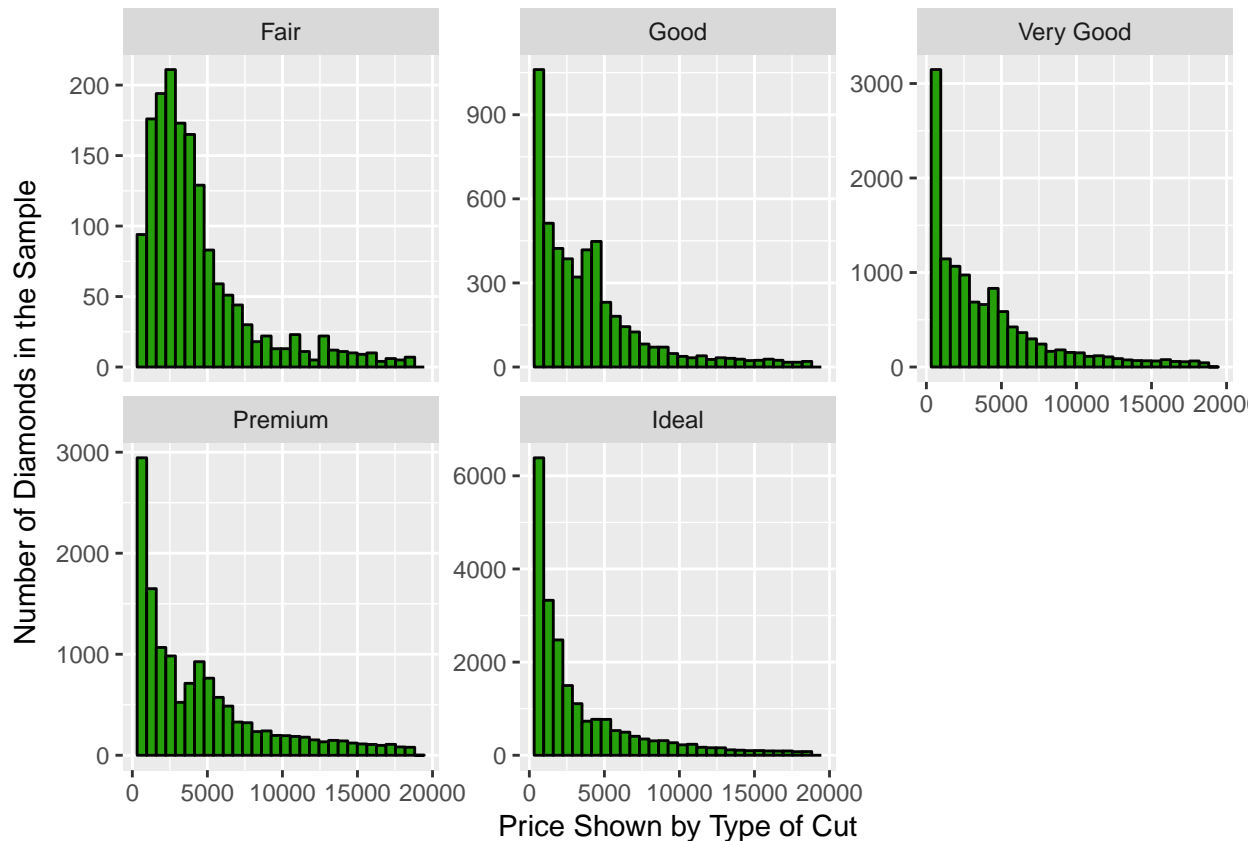
Scales and Multiple Histograms

Getting a summary of the prices based on cuts shows that the distribution of the prices should be somewhat the same. However, in the histograms I created for prices based on cuts, we saw that only the distributions of Premium, Very Good and Ideal look the same. The distribution of Fair and Good look somehow uniform.

Let's look into this more. I now create histograms of prices based on cuts, where the y-axis is not fixed. I do it by scaling the y-axis.

```
qplot(x = price, data = diamonds,
      color = I('black'), fill = I('#25A00A')) +
  xlab('Price Shown by Type of Cut') +
  ylab('Number of Diamonds in the Sample') +
  facet_wrap(~cut, scales="free_y")

## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```



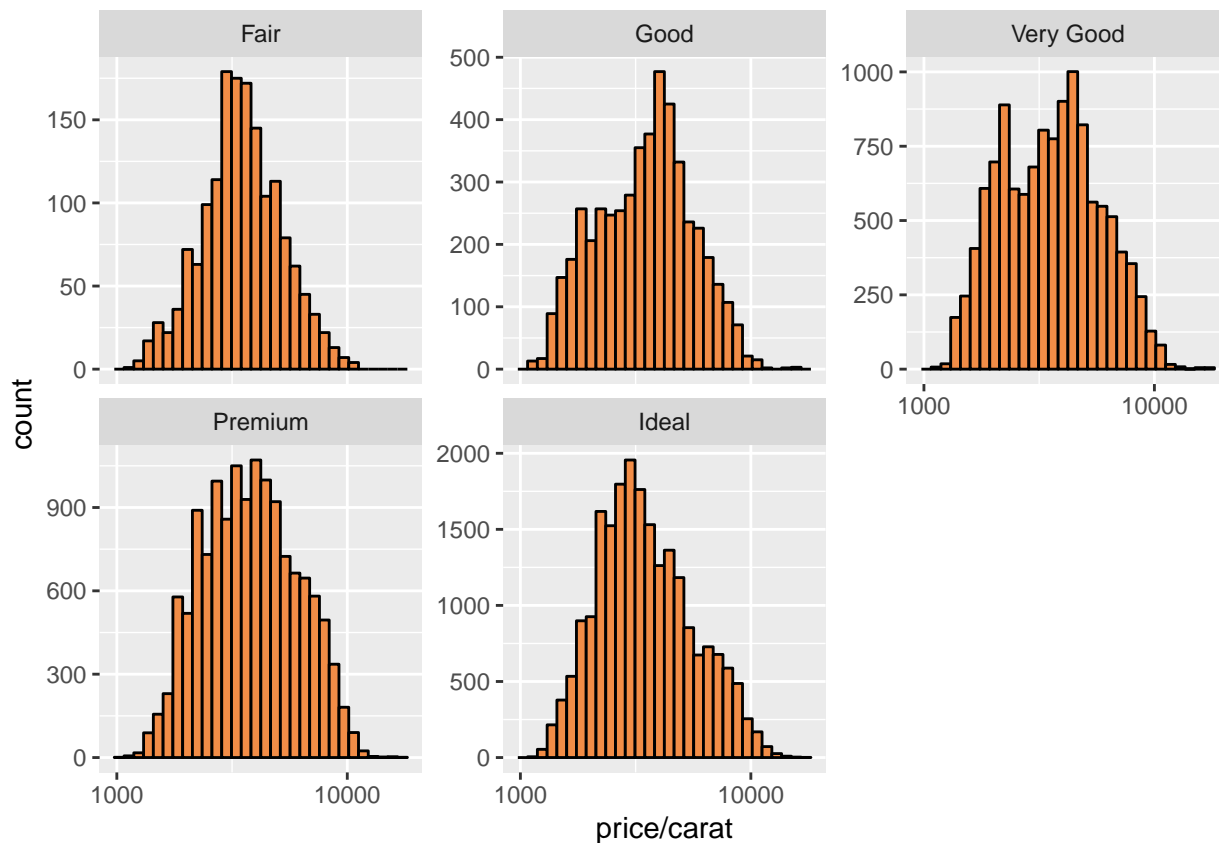
Scaling the y-axis makes it different on some of the panels in the histogram. For example, for Fair cuts, the y-axis goes from 0 to over 200, whereas, in Premium it goes from 0 to about 3000.

Price per Carat by Cut

I want to create a histogram of price per carat and facet it by cut. Since histograms only takes a single series and splits it into bins, for the 'x' variable, I use 'price/carat' to get the distribution for price per carat. I wrap it by cut, and add scales for the y-axis to get different values on the y-axis. I complete my histogram by scaling it using log10. After adding log10, I can see how the distribution has changed to a somehow normal one.

```
qplot(data = diamonds, x = price/carat,
      color = I('black'), fill = I('#F28D47')) +
  scale_x_log10() +
  facet_wrap(~cut, scales = "free_y")
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```



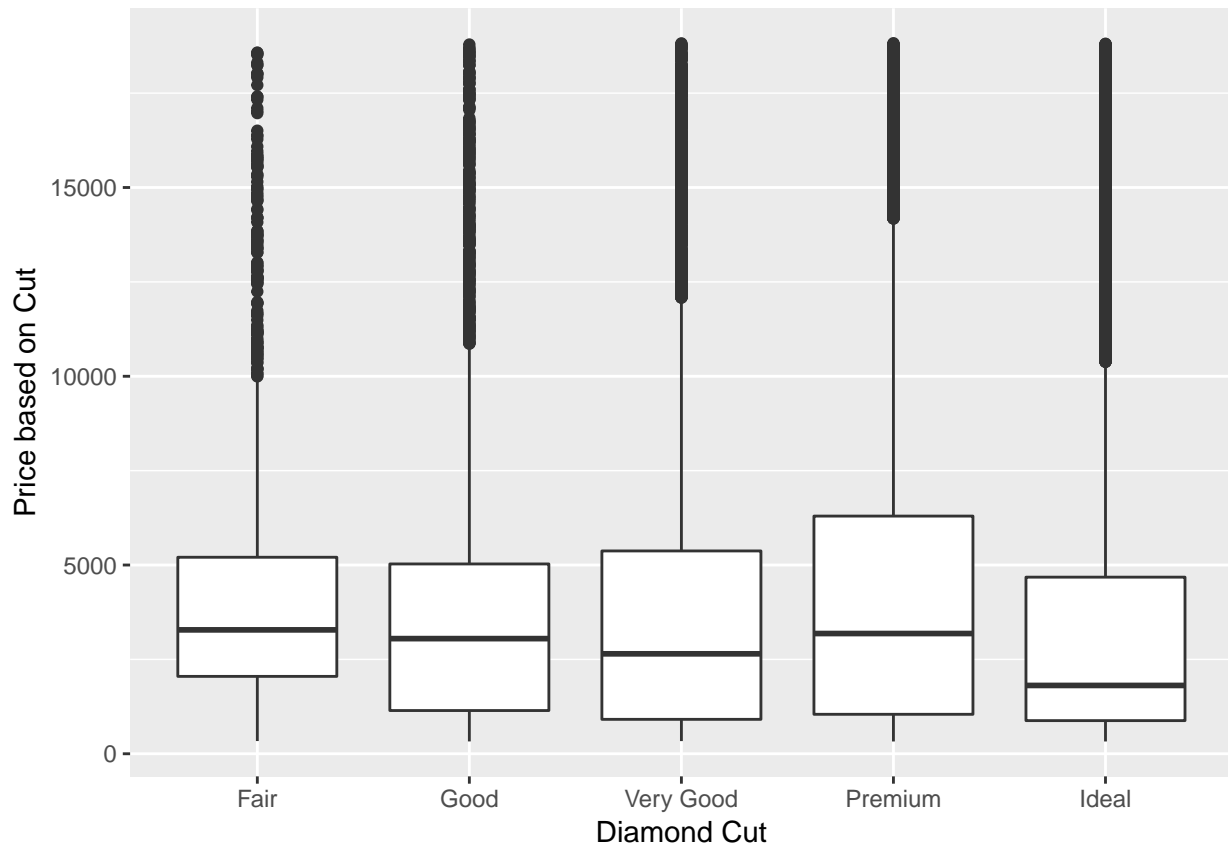
Price Box Plots

I want to investigate the price of diamonds using box plots, numerical summaries, and the following categorical variables: cut, clarity, or color.

Box Plot by Cut

To get the boxplot based on cut, I set the x-axis to show different cuts, and the y-axis to show the price of those cuts.

```
qplot(data = diamonds, geom = 'boxplot',
      x = cut, y = price) +
  xlab('Diamond Cut') +
  ylab('Price based on Cut')
```

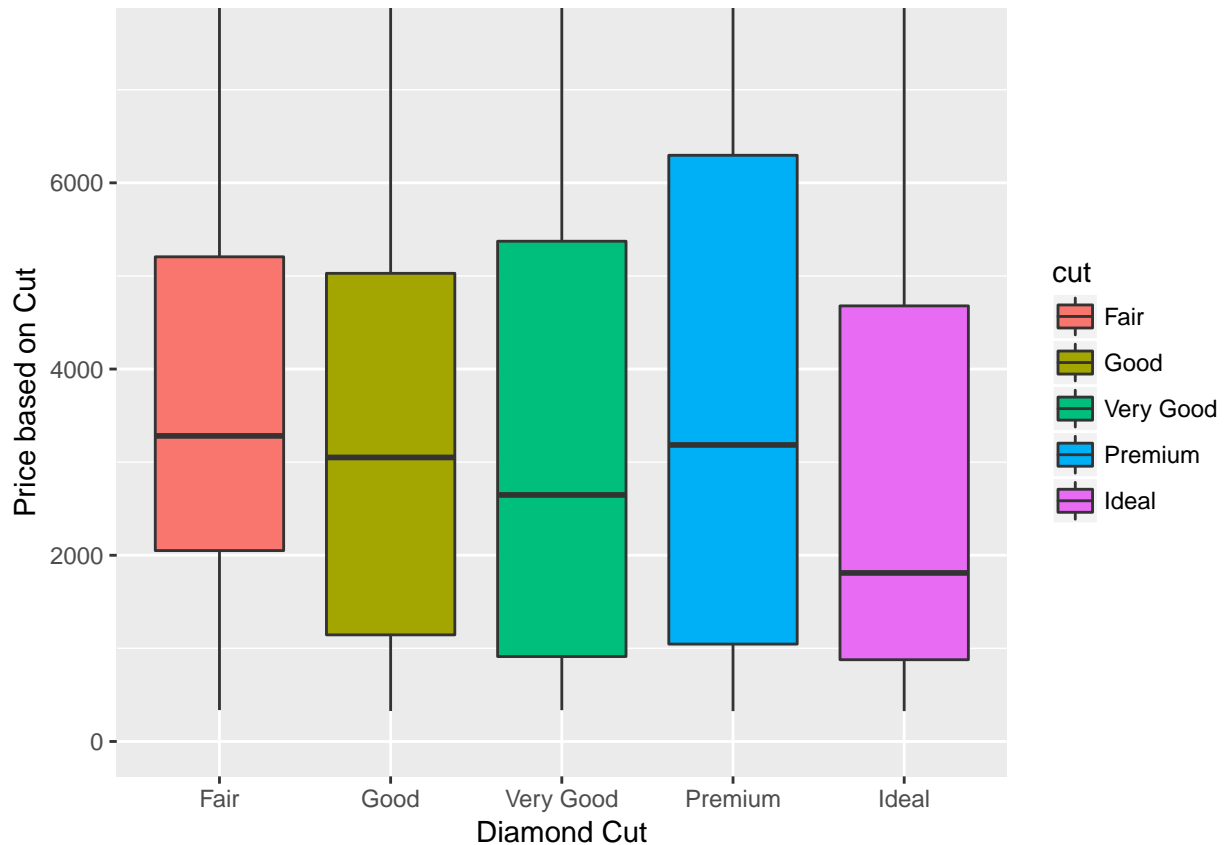



These boxplots show there are a lot of outliers which fall over \$10000 price. I need to get a closer look at the boxes without having to deal with the outliers. To do that, I use the 'coord_cartesian' which can magnify the boxes for me.

The Cartesian coordinate system is the most familiar, and common, type of coordinate system. Setting limits on the coordinate system will zoom the plot (like you're looking at it with a magnifying glass), and will not change the underlying data like setting limits on a scale will. Read more on: http://docs.ggplot2.org/0.9.3.1/coord_cartesian.html

Since the boxplots lie between 0 and \$7000, I adjust my y-axis to limit the prices up to 7500. I also add a bit of color to the boxplots for a nicer look.

```
qplot(data = diamonds, geom = 'boxplot',
      x = cut, y = price,
      fill=cut) +
  xlab('Diamond Cut') +
  ylab('Price based on Cut') +
  coord_cartesian(ylim = c(0, 7500))
```



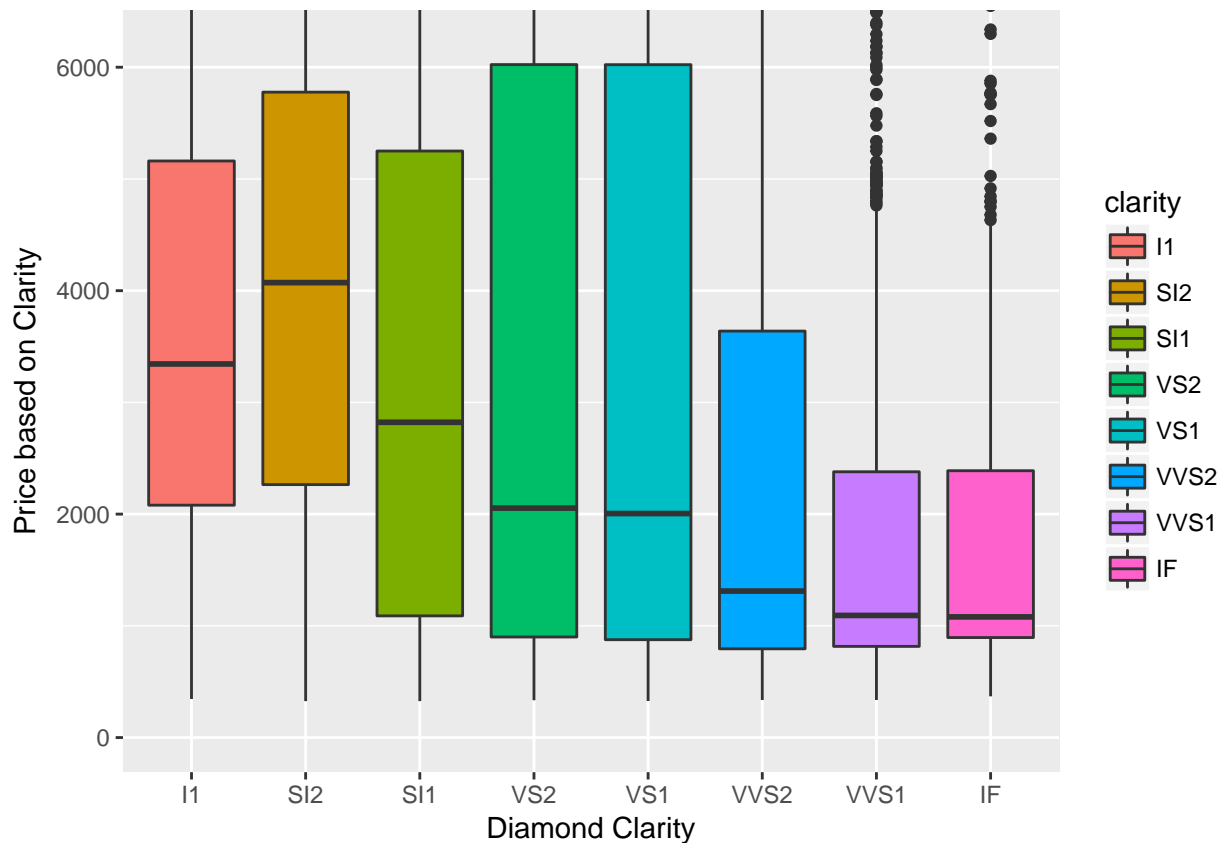
The barplots show that the Premium cut has the largest range and highest median (i.e. the black horizontal line in the box), where the Ideal cut has the lower price.

Box Plot by Clarity

To get the boxplot based on clarity, I set the x-axis to show different clarities, and the y-axis to show the price of them.

I first generate my boxplot without using cartesian coordinates. This will give me the chance to see the ranges and where the outliers are. After investigating these things, I adjust my coordinates.

```
qplot(data = diamonds, geom = 'boxplot',
      x = clarity, y = price,
      fill=clarity) +
  xlab('Diamond Clarity') +
  ylab('Price based on Clarity') +
  coord_cartesian(ylim = c(0, 6200))
```



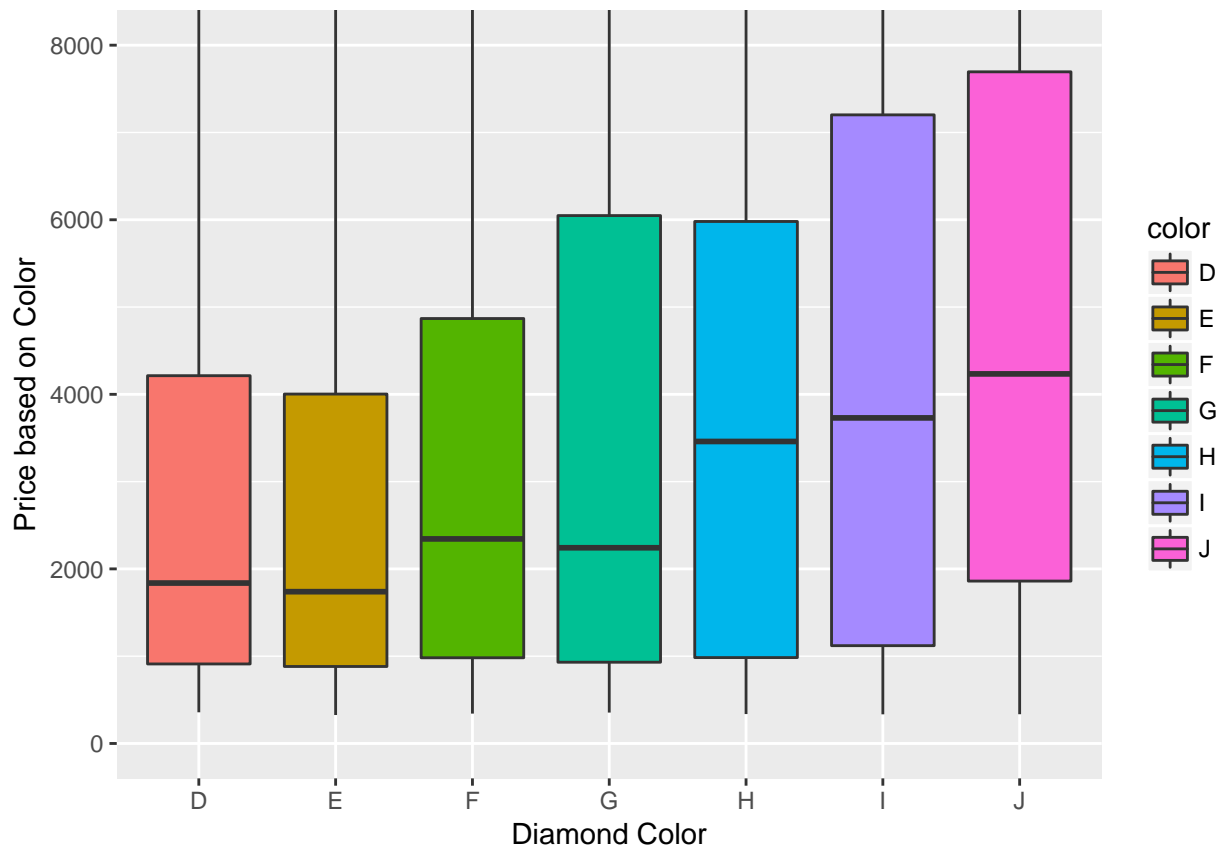
The barplot show VS1 and VS2 have somehow the same range, with the Median of VS2 being slightly larger than VS1. The largest Median in price goes to SI2 clarity category.

Box Plot by Color

To get the boxplot based on clarity, I set the x-axis to show different colors, and the y-axis to show the price of them.

I first generate my boxplot without using cartesian coordinates. This will give me the chance to see the ranges and where the outliers are. After investigating these things, I adjust my coordinates.

```
qplot(data = diamonds, geom = 'boxplot',
      x = color, y = price,
      fill=color) +
  xlab('Diamond Color') +
  ylab('Price based on Color') +
  coord_cartesian(ylim = c(0, 8000))
```



The barplots show that color E seems to have the lowest price, and the largest Median belongs to color J.

Interquartile Range - IQR

Now I will work with the `IQR()` function and try to get some info regarding price and color of diamonds. `IQR()` computes the interquartile range of the `x` values.

Let's first investigate a bit on the price range based on color. I use the 'by' command to get a summary of prices, categorized by color.

```
by(diamonds$price,diamonds$color, summary)
```

```
## diamonds$color: D
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##   357    911    1838    3170    4214   18690
## -----
## diamonds$color: E
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##   326    882    1739    3077    4003   18730
## -----
## diamonds$color: F
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##   342    982    2344    3725    4868   18790
## -----
## diamonds$color: G
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##   354    931    2242    3999    6048   18820
```

```
## -----
## diamonds$color: H
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##   337     984     3460    4487    5980    18800
## -----
## diamonds$color: I
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##   334     1120     3730    5092    7202    18820
## -----
## diamonds$color: J
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##   335     1860     4234    5324    7695    18710
```

Based on this summary, I can for example say that the first quartile(25%) of diamonds with color D has the price 911 dollars, and the third quartile(75%) of them are 4214 dollars.

What is the IQR for the diamonds with the worst color?

Based on the documentation for the diamonds dataset, in diamonds colors, 'J' is the worst and 'D' is the best.

```
IQR(subset(diamonds, color=="J")$price)
```

```
## [1] 5834.5
```

What is the IQR for the diamonds with the best color?

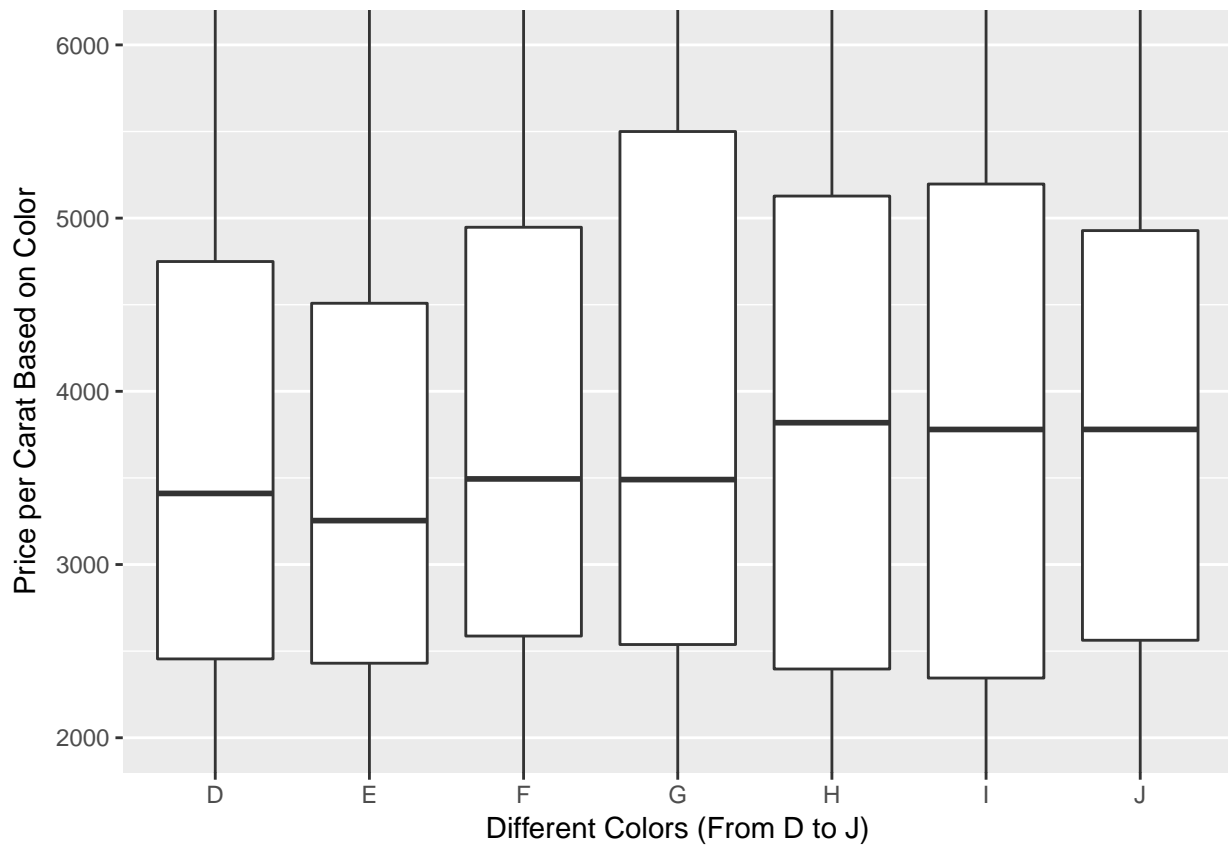
```
IQR(subset(diamonds, color=="D")$price)
```

```
## [1] 3302.5
```

Price per Carat Box Plots by Color

I want to investigate the price per carat of diamonds across the different colors of diamonds using boxplots.

```
qplot(data = diamonds, geom = 'boxplot',
      x = color, y = price/carat) +
  xlab('Different Colors (From D to J)') +
  ylab('Price per Carat Based on Color') +
  coord_cartesian(ylim = c(2000, 6000))
```

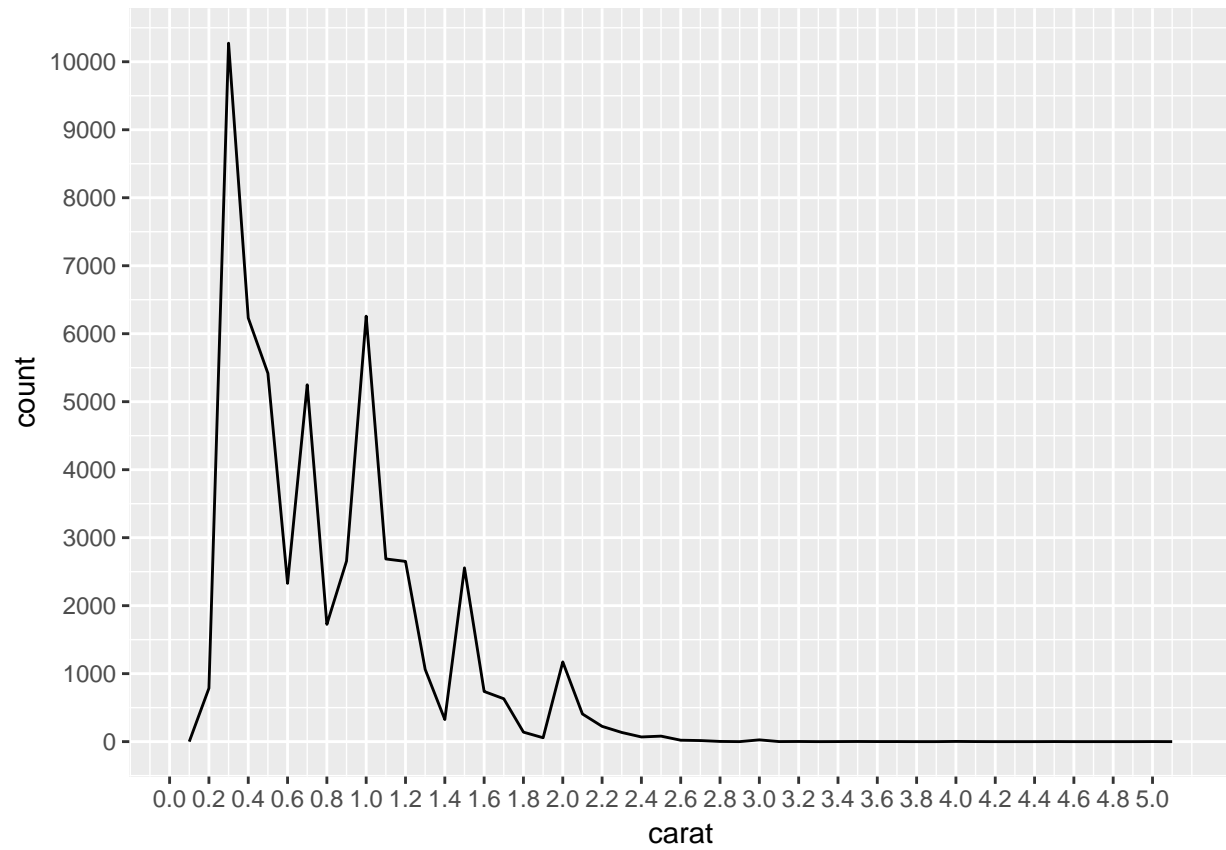


Among different colors, color 'G' has the most range, and color 'I' has the lowest price. The highest Median goes to color 'H'.

Carat Frequency Polygon

Using a frequency polygon, I want to find out what carat size has a count greater than 2000.

```
qplot(data = diamonds, x = carat, binwidth = 0.1,
      geom = 'freqpoly') +
  scale_y_continuous(breaks = seq(0, 15000, 1000)) +
  scale_x_continuous(breaks = seq(0, 5, 0.2))
```



From this plot, I can see that for example 0.3 and 1.01 sizes (carats) have more than 2000 counts, while anything after 2.0 has fewer than 2000 counts.